

Recurrent Neural Networks for Time Series Prediction

A little about me

I teach Data Science at [Galvanize](#).

Background

Ph.D. thesis: “Transition Regime Heat Conduction in Parabolic Trough Receivers”

Worked in concentrating solar thermal power for a decade (NREL and Abengoa)

Present interests:

I’ve been playing with neural nets for about a year. Useful resources:

[Andrej Karpathy’s NN course](#), [Karpathy Github](#)

[Ian Goodfellow’s Deep Learning book](#)

[Iamtrask’s blog](#)

[Christopher Olah’s blog](#)

[Jakob Aungiers’s blog](#), [Aungiers Github](#)

This presentation and source code:

https://github.com/GalvanizeOpenSource/Recurrent_Neural_Net_Meetup




Time series

A series of data points
(usually in time).

Data are partly dependent
on data that came before.

The order of data points
matters.



The quick brown fox jumps over the lazy dog.

The quick brown dog jumps over the lazy fox.

Recurrent Neural Networks (RNNs)


Models based on the connection of simple computational units.

Connections between units can form a directed cycle.
This gives a network the ability to maintain **a state based on previous inputs**.

This **state** helps it model time series.

*They pulled the boat up to the river **bank**.*

*They had been casing the **bank** for months.*



RNN use case - Handwriting generation

recurrent neural network handwriting generation demo

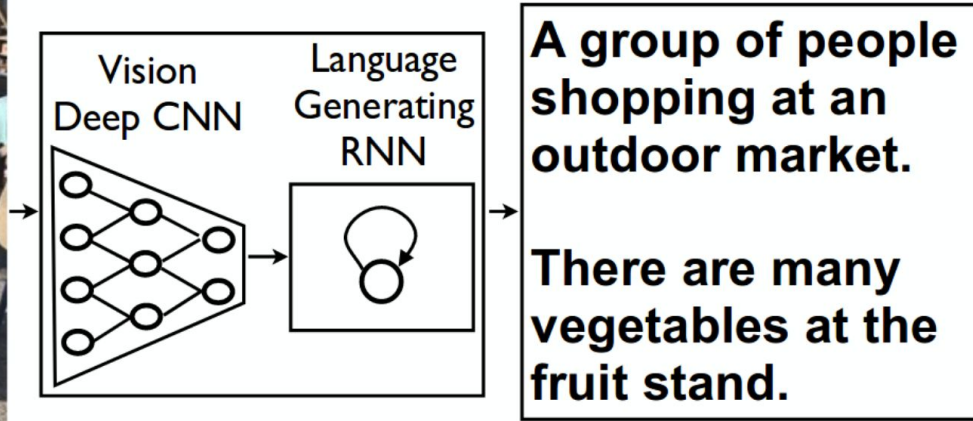
Type a message into the text box, and the network will try to write it out longhand ([this paper](#) explains how it works, source code is available [here](#)). Be patient, it can take a while!

Text --- up to 100 characters, lower case letters work best

The quick brown fox jumped over the lazy dog.

The quick brown fox jumped over the lazy dog.
The quick brown fox jumped over the lazy dog.

RNN use case - Image Captioning



RNN use case - predicting stock closing price



epochs = 1, window size = 50, sequence shift = 50

Mathematical & code understanding of RNNs

Start with a simple neural network (multilayer perceptron)

- Show how you would code a simple one in Python

Extrapolate from that to a simple recurrent neural network

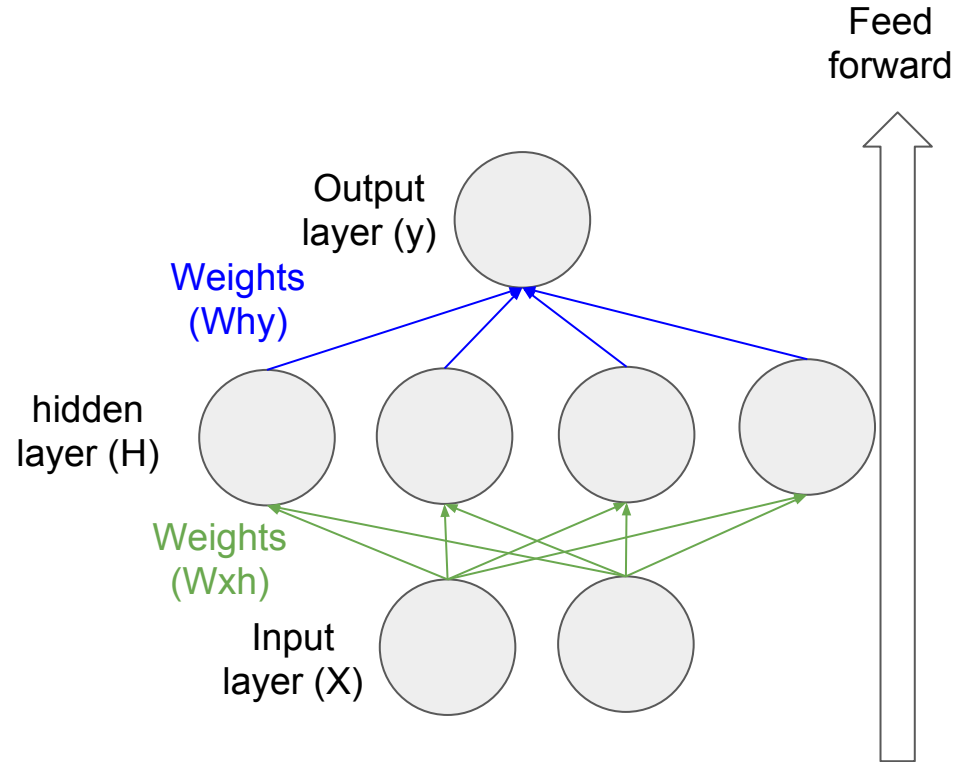
- Show how the code is modified from the MLP

Will use similar code to make a new Dr. Seuss book.

Finish up with LSTMs and stock prediction.


Multilayer perceptron

- Maps input data to corresponding outputs. Calculation *feeds forward* through the network.
- Nodes are arranged in layers.
- Nodes have values for any input given the sum of inputs to the node and an *activation function* that transforms the sum to a non-linear output.
- The “learning” in the network is held by the trained values of the connections (the weights). These weights start with random values.




Multilayer perceptron - learning the weights

- After feed forward yields a predicted output (\mathbf{y}_p), its difference (loss) is calculated from the real output (\mathbf{y}).
- Back propagation estimates how much the loss varies due to each weight in the network (the gradient).
- Gradient descent uses the gradient and learning rate to tweak each of the weights so that the network predicts a little better next time.
- When the total loss reaches an acceptable level, stop. Now it's trained and ready to predict.




MLP code




```
31 # initialize weights
32 Wxh = 2*np.random.uniform(size=(nodes_input, nodes_hidden)) - 1
33 Why = 2*np.random.uniform(size=(nodes_hidden, nodes_target)) - 1
34
35 # simulation parameters
36 np.random.seed(1)
37 alpha = 1 # learning rate
38 num_epochs = 20000 # number of epochs
39
40 # training
41 print("\nTraining:")
42 for e in range(num_epochs):
43     yp_lst = [] # predictions
44     error_lst = [] # differences between target and predictions
45     for X, y in zip(inputs, targets):
46         X = X.reshape((1, X.shape[0])) # for row, column shape
47         # Feed forward
48         H = activation(np.dot(X,Wxh))
49         yp = activation(np.dot(H,Why))
50         # Back propogate to find gradients
51         # Why gradients
52         yp_error = y - yp
53         yp_delta = yp_error*activation(yp,deriv=True)
54         grad_Why = np.dot(H.T, yp_delta)
55         # Wxh gradients
56         H_error = np.dot(yp_delta, Why.T)
57         H_delta = H_error * activation(H,deriv=True)
58         grad_Wxh = np.dot(X.T, H_delta)
59         # Use gradient descent to update weights
60         Why += alpha * grad_Why
61         Wxh += alpha * grad_Wxh
```

Comparing the simplest version of these neural nets

MLP




RNN



A double arrow indicates a weight in each direction (2 weights).


So how many weights in each architecture?

MLP



Why =
W_{xh} =


RNN



Why =
W_{hh} =
W_{xh} =


So how many weights in each architecture?

MLP



$$W_{hy} = 4$$
$$W_{xh} = 8$$


RNN




$$W_{hy} = 4$$
$$W_{hh} = 16$$
$$W_{xh} = 8$$

Benefit of the intra - layer recurrent connections

- The previous state of a node in a recurrent hidden layer (**H_{prev}** for coding purposes) can affect the value of itself or other nodes in the layer.
- This gives the net the ability to model sequential data.
- Feedforward and backpropagation work the same way.
- Learn **W_{hh}** like all the other weights. In a trained model all the weights are fixed. It's the activations of the nodes that changes with changes in sequence.



RNN code




https://github.com/GalvanizeOpenSource/Recurrent_Neural_Net_Meetup/blob/master/rnn_soln.py

```
96 # training
97 print("\nTraining:")
98 H_prev = np.zeros((1, nodes_hidden))
99 H_delta_fut = np.zeros(nodes_hidden)
100 for e in range(num_epochs):
101     error_lst = [] # differences between target and predictions
102     for X, y in zip(X_train, y_train):
103         X = X.reshape((1, X.shape[0])) # for row, column shape
104         y = y.reshape((1,1))
105         # Feed forward
106         H = activation(np.dot(X,Wxh) + np.dot(H_prev,Whh))
107         yp = activation(np.dot(H,Why))
108         # Back propogate to find gradients
109         # Why gradients
110         yp_error = y - yp
111         yp_delta = yp_error*activation(yp,deriv=True)
112         grad_Why = np.dot(H.T, yp_delta)
113         # Wxh gradients
114         H_error = np.dot(yp_delta, Why.T) + np.dot(H_delta_fut, Whh.T)
115         H_delta = H_error * activation(H,deriv=True)
116         #H_delta_fut = np.copy(H_delta) crashes simulation
117         grad_Wxh = np.dot(X.T, H_delta)
118         # Whh gradients
119         grad_Whh = np.dot(H_prev.T, H_delta)
120         # Use gradient descent to update weights
121         Why += alpha * grad_Why
122         Whh += alpha * grad_Whh
123         Wxh += alpha * grad_Wxh
124         # save for future use
125         H_prev = np.copy(H)
```





RNN - text is sequential data

Use the `min-char-rnn.py` code to learn Dr. Seuss. As the model trains it will eventually write some new books!



Moving into multilayer RNNs

Multiple layers (and more nodes in each layer) allow more difficult sequences to be learned. They are also harder to train. Exploding and vanishing gradients cause convergence problems, too.



Keras

[Keras](#) is a high-level neural networks API, written in Python and capable of running on top of either TensorFlow, CNTK or Theano.

We use it in the DSI for capstone projects. MLPs, CNNs, RNNs, some Reinforcement Learning too.

[An API for TensorFlow](#)


Available recurrent layers:

- Recurrent

- SimpleRNN

- Long Short-Term Memory (LSTM)

- Gated Recurrent Unit (GRU)

 Keras Documentation

Search docs

[Home](#)
[Getting started](#)

- [Guide to the Sequential model](#)
- [Guide to the Functional API](#)
- [FAQ](#)

[Models](#)

- [About Keras models](#)
- [Sequential](#)
- [Model \(functional API\)](#)

[Layers](#)


- [About Keras layers](#)
- [Core Layers](#)
- [Convolutional Layers](#)
- [Pooling Layers](#)
- [Locally-connected Layers](#)

[Recurrent Layers](#)

- [Recurrent](#)
- [SimpleRNN](#)
- [GRU](#)
- [LSTM](#)

LSTM

- Long short-term memory (LSTM) is an architecture (an artificial neural network) proposed in 1997.
- LSTM network is well-suited ... when there are time lags of unknown size and bound between important events.
- LSTM practical applications: natural language text compression, handwriting recognition, speech recognition, translation.
([See Wikipedia](#))
- Addresses the vanishing gradient problem.
- [Christopher Olah's blog](#) to gain insight.

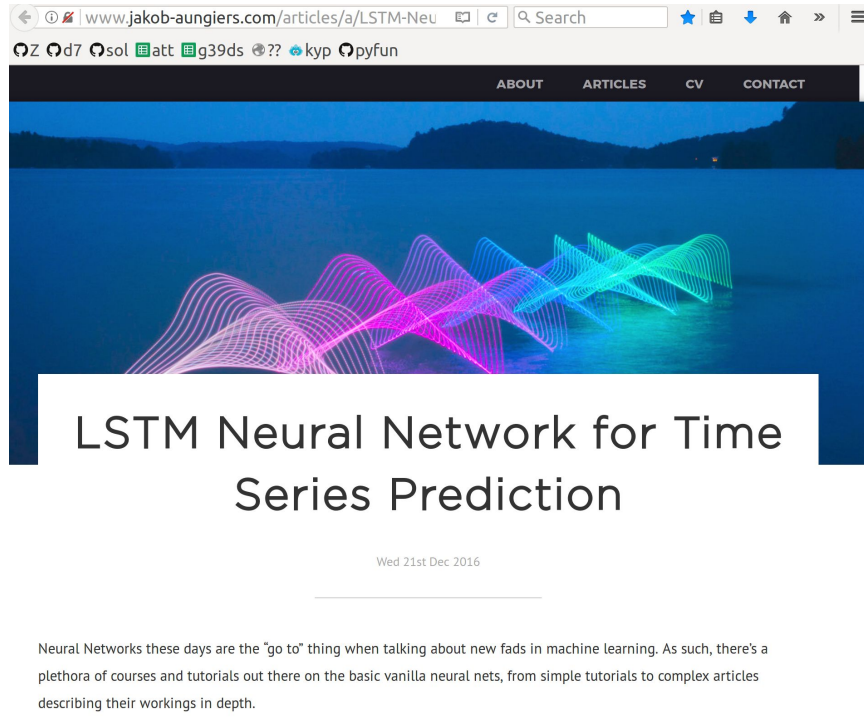


[Attribute](#)

Using Keras LSTM layers to predict stock price


See `lstm.py` and `run_lstm.py` in the Github repo. These files were written by Jakob Aungiers. Check out his [blog](#) and [Github](#).

Jakob uses 2 LSTM layers to predict the S&P500.



Code and results

```
46 def build_model(layers):
47     model = Sequential()
48
49     model.add(LSTM(
50         input_shape=(layers[1], layers[0]),
51         output_dim=layers[1],
52         return_sequences=True))
53     model.add(Activation("tanh"))
54     model.add(Dropout(0.2))
55
56     model.add(LSTM(
57         layers[2],
58         return_sequences=False))
59     #model.add(Activation("tanh"))
60     model.add(Dropout(0.2))
61
62     model.add(Dense(
63         output_dim=layers[3]))
64     model.add(Activation("linear"))
65
66     start = time.time()
67     model.compile(loss="mse", optimizer="rmsprop")
68     print("> Compilation Time : ", time.time() - start)
69     return model
```



https://github.com/GalvanizeOpenSource/Recurrent_Neural_Net_Meetup/blob/master/lstm.py

Thank you!

You can contact me at frank.burkholder@galvanize.com